

# 1. Introduction

This report aims to provide a thorough analysis of the Spark application our group has designed and implemented.

During the process of designing and implementing the application, PySpark has been used as the programming language to perform the exploratory data analysis. For the purpose of achieving the most efficient analysis outcomes, the AWS EMR spark cluster has been chosen as the running environment to support our spark application. The program has been processed on a cluster of a single **r5.4xlarge** node.

In order to perform the distributed computation using Pyspark, Spark RDD, Spark dataframe and Spark ML have been used to perform a series of operations. The distributed file system Amazon S3 has been utilised to store the data.

## 2. Vocabulary Exploration

### 2.1 Final Results

*Results 1:*

Requirements	Results
Common words between <b>matched</b> & <b>mismatched</b> sets	8927
The number of words <b>unique</b> to the <b>matched</b> sets	11134
The number of words unique to <b>mismatched</b> sets	6466

*Results 2: Percentage of words appearing*

Frequency	Percentage of appearing (contains stop words)	Percentage of appearing (without stop words)
1	69.5084%	69.6106%
2	11.0784%	11.0901%
3	6.4515%	6.4593%
4	5.2362%	5.2389%
5	7.7255%	7.601%

## 2.2 Workflow

### Part 1 Workflow

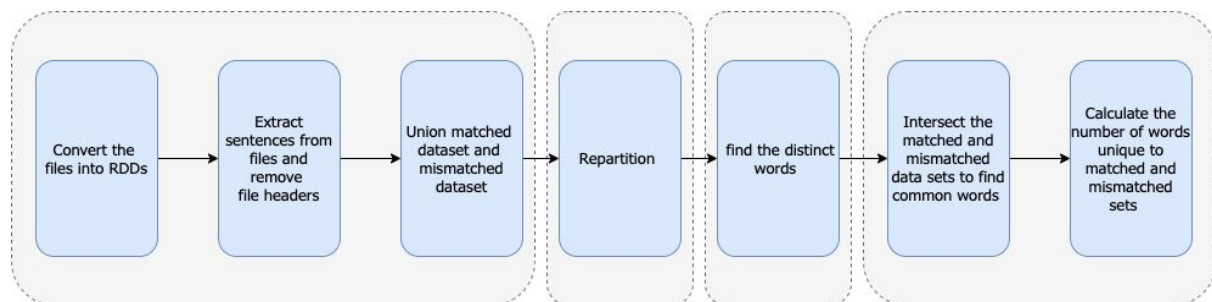


Figure 1. Count the matched & mismatched words across four data sets

The workflow of counting the matched and mismatched words across four different datasets: dev\_matched.tsv, dev\_mismatched.tsv, test\_matched.tsv and test\_mismatched.tsv has been concluded in figure 1. It can be seen that four shuffle transformations were involved in this process, invoked by the actions of Union(), repartition(), distinct() and intersection() respectively.

### Part 2 Workflow

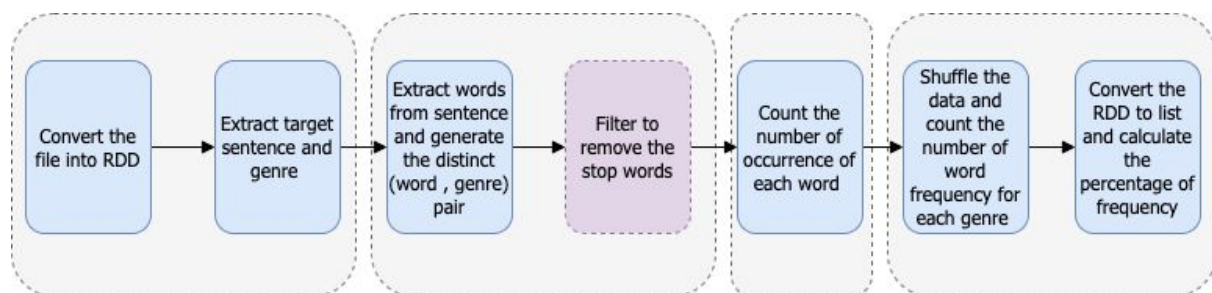


Figure 2. Calculate the percentage of words appearing

The process of calculating the percentage of words appearing frequencies has been included in figure2 & figure3. It can be seen that there were three shuffle transformations was involved in this process:

- ❑ Generate the distinct (word, genre) pair (by the distinct() action)
- ❑ Count the occurrence of each word under a particular genre (by the reduceByKey() action)
- ❑ Count the number of occurrence of each genre (by the reduceByKey() action)

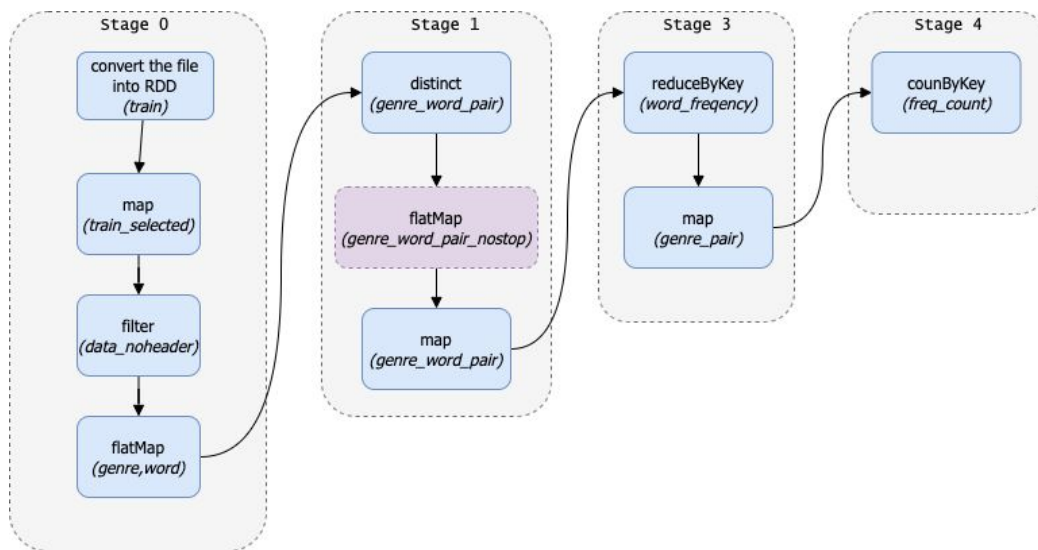


Figure 3. Data flow DAG

## 2.3 Implemented features

Various programming methodologies have been implemented to improve the permanence of the overall data analytics process.

### Repartition

Progress for count at <stdin>:46		Job Progress: 88/88 Tasks Complete		
Stage [ID]: name at [source]:[line]	Status	Task Progress	Elapsed Time (seconds)	Failed Task Logs
Stage [91]: coalesce at Nati...java:0	COMPLETE	4/4	6.83	
Stage [92]: distinct at <stdin>:42	COMPLETE	16/16	1.225	
Stage [93]: distinct at <stdin>:41	COMPLETE	16/16	1.541	
Stage [94]: intersection at ...in>:45	COMPLETE	32/32	0.909	
Stage [95]: count at <stdin>:46	COMPLETE	16/16	0.25	

#### ▼ Job [66]: count at <stdin>:46

Progress for count at <stdin>:46		Job Progress: 56/56 Tasks Complete		
Stage [ID]: name at [source]:[line]	Status	Task Progress	Elapsed Time (seconds)	Failed Task Logs
Stage [122]: distinct at <stdin>:42	COMPLETE	4/4	7.341	
Stage [123]: distinct at <stdin>:41	COMPLETE	4/4	7.094	
Stage [124]: intersection at ...in>:45	COMPLETE	32/32	0.882	
Stage [125]: count at <stdin>:46	COMPLETE	16/16	0.267	

Figure 4. Execution statistics with(first) / without(second) the action of repartition()

It can be seen from figure 4, the action of repartition() significantly improved the program performance by reducing the execution time from 15.584 seconds to 4.006 seconds.

## Closure

The stop words list variable has been broadcasted to every node using efficient broadcast algorithms to reduce communication costs. In our case, the stop words list has been broadcasted to every node. The session below shows this process in details:

```
1 stop_words = stopwords.words('english')
2 stop_words = sorted(stop_words)
3 broadcast_stopwords = sc.broadcast(stop_words)
```

Figure 5. Create the broadcast variable

```
1 stopwords_value = broadcast_stopwords.value
```

Figure 6. Receive the broadcast variable values

## Binary search on stop words

The binary search technique has been utilised to remove the stop words from the given list of words. Before performing the binary search, the stop words list has been sorted in ascending order. The flow chart of the searching process is shown in figure 7.

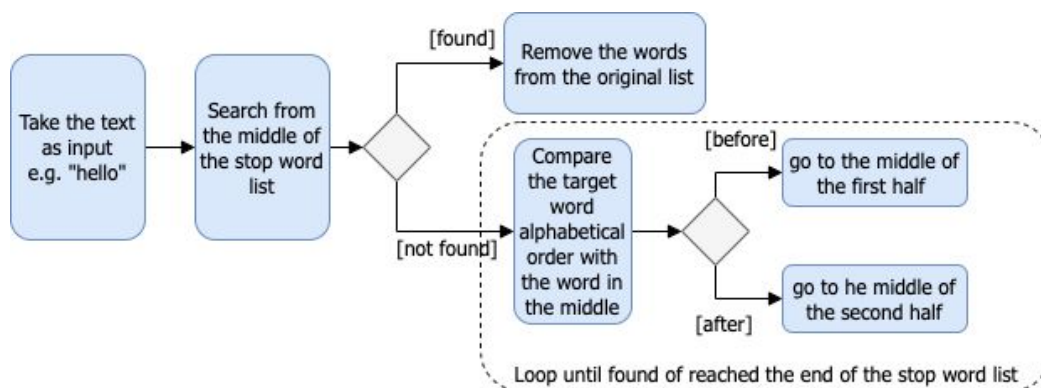


Figure 7. Binary search on stop words workflow

# 4. Sentence Vector Exploration

## 4.1 Method 1: TF-IDF Based Vector Representation

Term frequency-inverse document frequency (TF-IDF) feature vectorization method has been implemented to generate the sentence features.

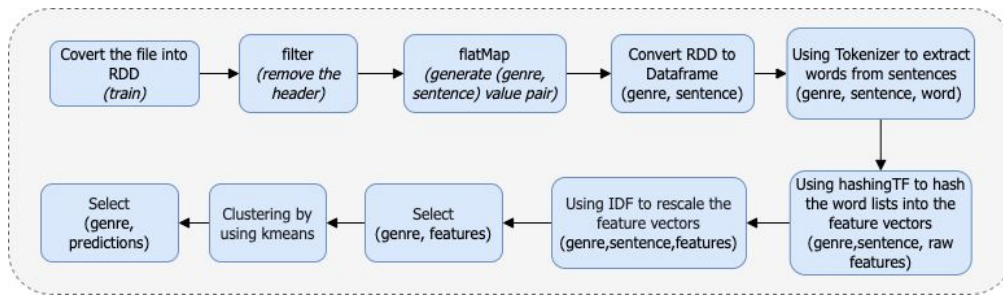


Figure 8. The workflow of the TF-IDF Based Vector Presentation process

During the sentence vector exploration process, the (sentence, genre) pair has been extracted from the given dataset. Then, Tokenizer was used to split sentences into words. For each sentence, HashingTF has been used to hash the list of words into a feature vector. Finally, IDF has been utilised to rescale the feature vectors, which improved the performance when using text as features.

## 4.2 Method 2: Pre-trained Sentence Encoder

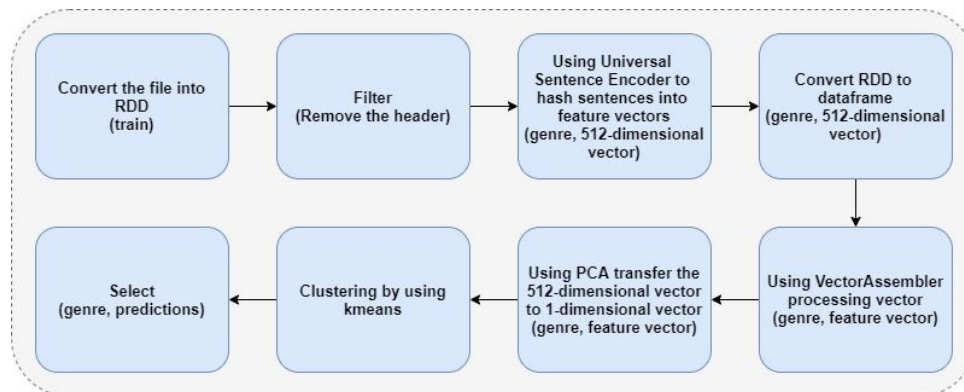


Figure 9. The workflow of the Pre-trained Sentence Encoder process

The encoding process by using the pre-trained sentence encoder has been shown in figure 9. Firstly, the encoder hashed the sentences into the raw feature vectors, which consisted of 512-dimensional vectors. Then, the VectorAssembler processing vector and Principal component analysis (PCA) have been used to convert the hashing results into a one-dimensional vector; prior to this process, the original RDD has been converted to a dataframe.

## 4.3 Clustering

Spark Machine learning K-means algorithm has been adopted as a clustering approach in order to divide the feature vectors into distinct groups based on some common characteristics. The K number has been set to 5, which means that the cluster has divided the input features into five prediction groups.

```

1 kmeans = KMeans().setK(5).setSeed(1)
2 model = kmeans.fit(feature)

```

Figure 10. K-means clustering

## Matrix 1 - TF-IDF Based Vector

Based on the clustering results, two confusion matrices have been presented for better visualisation of the results.

Prediction/True	fiction	government	slate	telephone	travel
0	88.33%	65.03%	76.15%	68.51%	65.01%
1	4.84%	7.16%	7.1%	2.49%	7.33%
2	0%	0%	0%	4.13%	0%
3	4.85%	27.62%	15.89%	2.07%	27.59%
4	0.98%	0.19%	0.02%	22.8%	0.07%

## Matrix 2 - Pre-trained Sentence Encoder

Prediction/True	fiction	telephone	travel	government	slate
0	46.17%	4.86%	3.89%	3.58%	13.86%
1	14.37%	56.58%	1.04%	1.79%	8.52%
2	10.66%	2.91%	50.71%	30.47%	27.52%
3	2.85%	7.06%	37.77%	50.89%	23.24%
4	25.95%	28.59%	6.59%	13.26%	26.86%

The confusion matrix above described the performance of the Kmeans cluster on a set of data for which the true values are known. The number of correct and incorrect predictions are summarised with count values and classified into five different prediction groups. The true positive predictions have been highlighted in grey.

By comparing matrix 1 and matrix 2, it can be seen that the prediction in matrix 2 is more accurate than it's in matrix 1. Based on the results from matrix 1, all genre types were classified into the same prediction group, conversely, each prediction group had a unique match with the true value (genre) in matrix 2. In a nutshell, the pre-trained sentence encoder performs better than the TF-IDF based vector in this case.

## 5. Performance Evaluation

### Performance Evaluation on Vocabulary Exploration Part 2

The table below summarises the differences between the two execution environments:

Settings/Environment	Environment 1	Environment 2
Execution Time In total	72.4 s	42s
Number of executor used	2	4
Number of threads	40	136
Total shuffling size	35M	31.8M

#### *Execution Time (Each stage)*

It can be seen that the application's execution time under environment 2 is approximately one time faster than when it's under environment 1. There were four stages involved in the application execution process; In stage 0, a text file has been converted to a resilient distributed dataset(RDD), it can be seen that it has taken approximately ten seconds under both environment 1 and environment 2. (Please see details in Appendix 3 & Appendix 8). Stage 1 of the job involved extracting the distinct key-value pair from the input data. Based on the execution statistics, it can be seen that the application under environment 2 processed two times faster than when it was under environment 1. Stage 2 of the job was invoked by the execution of `ReduceByKey()`, it can be seen that the application had a similar execution time under both environments, the major reason was that the data was reshuffled/repartitioned in stage two through the `distinct()` execution, data has the similar structure has been partitioned into one node, hence, this has reduced the overall execution time. Similarly, the execution time in stage 3 which was initiated by the `countByKey()` action had a similar execution time under both environments.

#### *Number of executors & threads*

There were four executors used under environment 2, which was twice of the executors used under environment 1. The same amount of jobs has been divided into a larger number of nodes, thus, a faster application execution time. (please see more details in Appendix 2 & Appendix 6).

## 6. Appendix

### Appendix 1. Cluster setup of environment 1 :

Root device EBS volume size  GiB ⓘ

Choose the instance type, number of instances, and a purchasing option. You can choose to use On-Demand Instances, Spot Instances, or both. The instance type and purchasing option apply to all EC2 instances in each instance group, and you can only specify these options for an instance group when you create it. [Learn more about instance purchasing options](#) ⓘ

Node type	Instance type	Instance count	Purchasing option
<b>Master</b> Master - 1 ⓘ	<b>m4.xlarge</b> ⓘ 8 vCore, 16 GiB memory, EBS only storage EBS Storage: 64 GiB ⓘ Add configuration settings ⓘ	1 Instances	<input checked="" type="radio"/> <b>On-demand</b> ⓘ <input type="radio"/> <b>Spot</b> ⓘ Use on-demand as max price ⓘ
<b>Core</b> Core - 2 ⓘ	<b>m4.xlarge</b> ⓘ 8 vCore, 16 GiB memory, EBS only storage EBS Storage: 32 GiB ⓘ Add configuration settings ⓘ	<input type="text" value="0"/> Instances	<input checked="" type="radio"/> <b>On-demand</b> ⓘ <input type="radio"/> <b>Spot</b> ⓘ Use on-demand as max price ⓘ

The program has been processed on a cluster of one m4.xlarge node (one master node) .

### Appendix 2. Execution statistics of environment 1 :

▼ Spark Job Progress

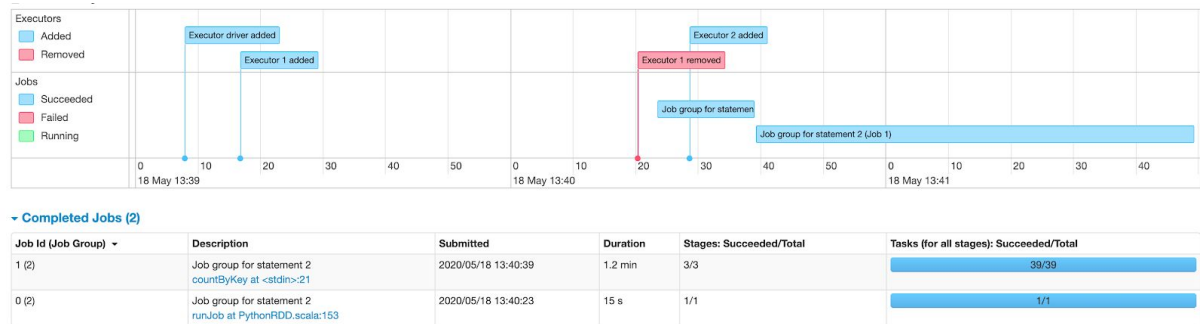
▼ Job [0]: runJob at PythonRDD.scala:153				
Progress for runJob at PythonRDD.scala:153		Job Progress: 1/1 Tasks Complete		
Stage [ID]: name at [source]:[line]	Status	Task Progress	Elapsed Time (seconds)	Failed Task Logs
Stage [0]: runJob at Python...la:153	COMPLETE	1/1	15.388	

▼ Job [1]: countByKey at <stdin>:21				
Progress for countByKey at <stdin>:21		Job Progress: 39/39 Tasks Complete		
Stage [ID]: name at [source]:[line]	Status	Task Progress	Elapsed Time (seconds)	Failed Task Logs
Stage [1]: distinct at <stdin>:14	COMPLETE	7/7	67.817	
Stage [2]: reduceByKey at <stdin>:17	COMPLETE	16/16	1.48	
Stage [3]: countByKey at <stdin>:21	COMPLETE	16/16	0.766	

Stage 1 had the heaviest workload. In this stage, the distinct key-value pair has been found.



### Appendix 3. Execution timeline of environment 1 :



The total execution time was 1.35 minutes and two executors have been used to perform the tasks.

### Appendix 4. Cluster setup of environment 2 :

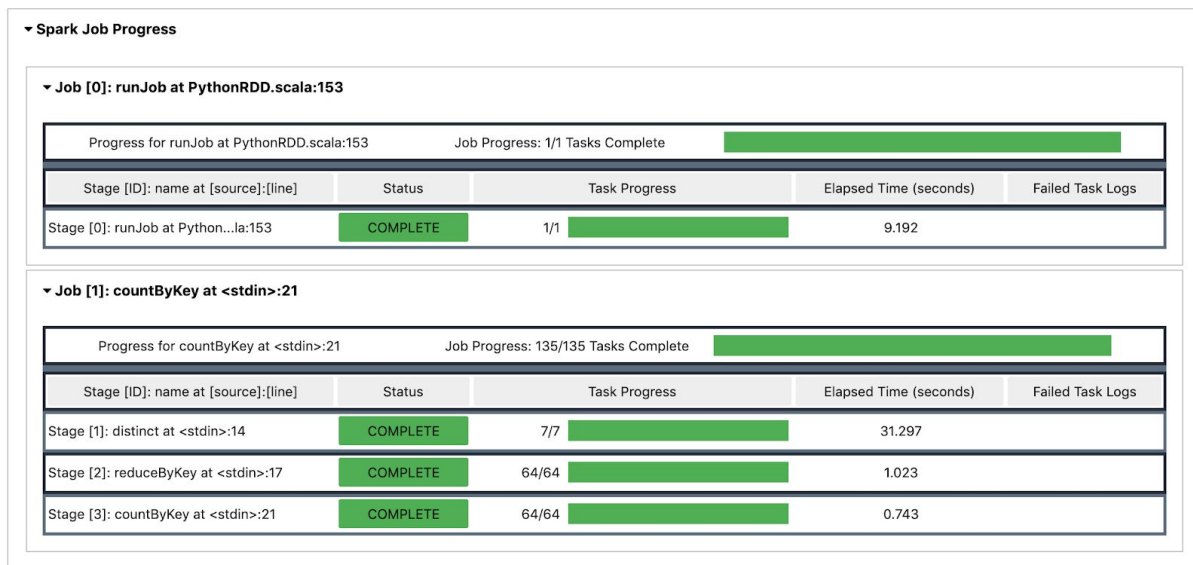
Root device EBS volume size  GiB ⓘ

Choose the instance type, number of instances, and a purchasing option. You can choose to use On-Demand Instances, Spot Instances, or both. The instance type and purchasing option apply to all EC2 instances in each instance group, and you can only specify these options for an instance group when you create it. [Learn more about instance purchasing options](#) ⓘ

Node type	Instance type	Instance count	Purchasing option
Master Master - 1 ⓘ	m4.xlarge ⓘ 8 vCore, 16 GiB memory, EBS only storage EBS Storage: 64 GiB ⓘ Add configuration settings ⓘ	1 Instances	<input checked="" type="radio"/> On-demand ⓘ <input type="radio"/> Spot ⓘ Use on-demand as max price ⓘ
Core Core - 2 ⓘ	m4.xlarge ⓘ 8 vCore, 16 GiB memory, EBS only storage EBS Storage: 64 GiB ⓘ Add configuration settings ⓘ	<input type="text" value="4"/> Instances	<input checked="" type="radio"/> On-demand ⓘ <input type="radio"/> Spot ⓘ Use on-demand as max price ⓘ


The program has been processed on a cluster of five m4.xlarge nodes (one master node and four slave nodes).

## Appendix 5. Execution statistics of environment 2 :



Stage 1 had the heaviest workload. In this stage, the distinct key-value pair has been found.

## Appendix 6. Execution timeline of environment 2 :



The Total execution time was 42 seconds and four executors have been used to perform the tasks.